


# Recap: Last Lecture

- Neural Networks
- Multilayer Neural Networks
- Backpropagation



Today:

# Convolutional Neural Networks

- CNN Basics
- Typical CNN Architectures


# Why CNN for Image?

100\*100\*3



Represented  
as pixels

$\mathbb{R}^{30000}$



The most basic  
classifiers


Use 1<sup>st</sup> layer as module  
to build classifiers

Use 2<sup>nd</sup> layer as  
module .....

Can the network be simplified by  
considering the properties of images?

# Why CNN for Image


- Hierarchical structure of objects.
  - Objects consist of object parts.
  - Object parts consist of simple, local patterns.



Low level features



Mid level features



High level features

# Why CNN for Image

- Some patterns are much smaller than the whole image

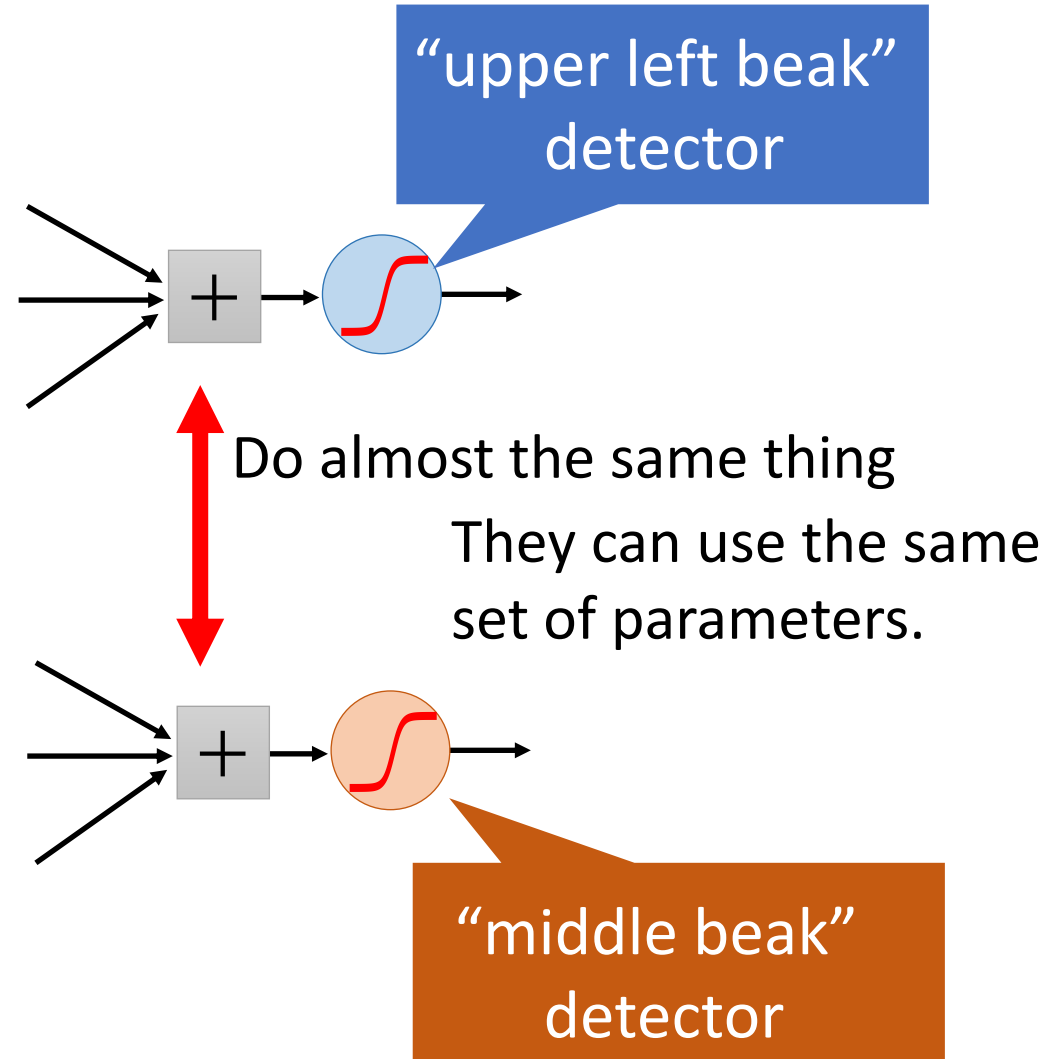

A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters



# Why CNN for Image

- Same objects or same patterns appear at different places of images.



# Why CNN for Image

- Subsampling/scaling the pixels will not change the object category

bird



subsampling

bird





We can subsample the pixels to make image smaller



Less parameters for the network to process the image

# The whole CNN

dog, cat, horse .....



Convolution

Max Pooling

Convolution

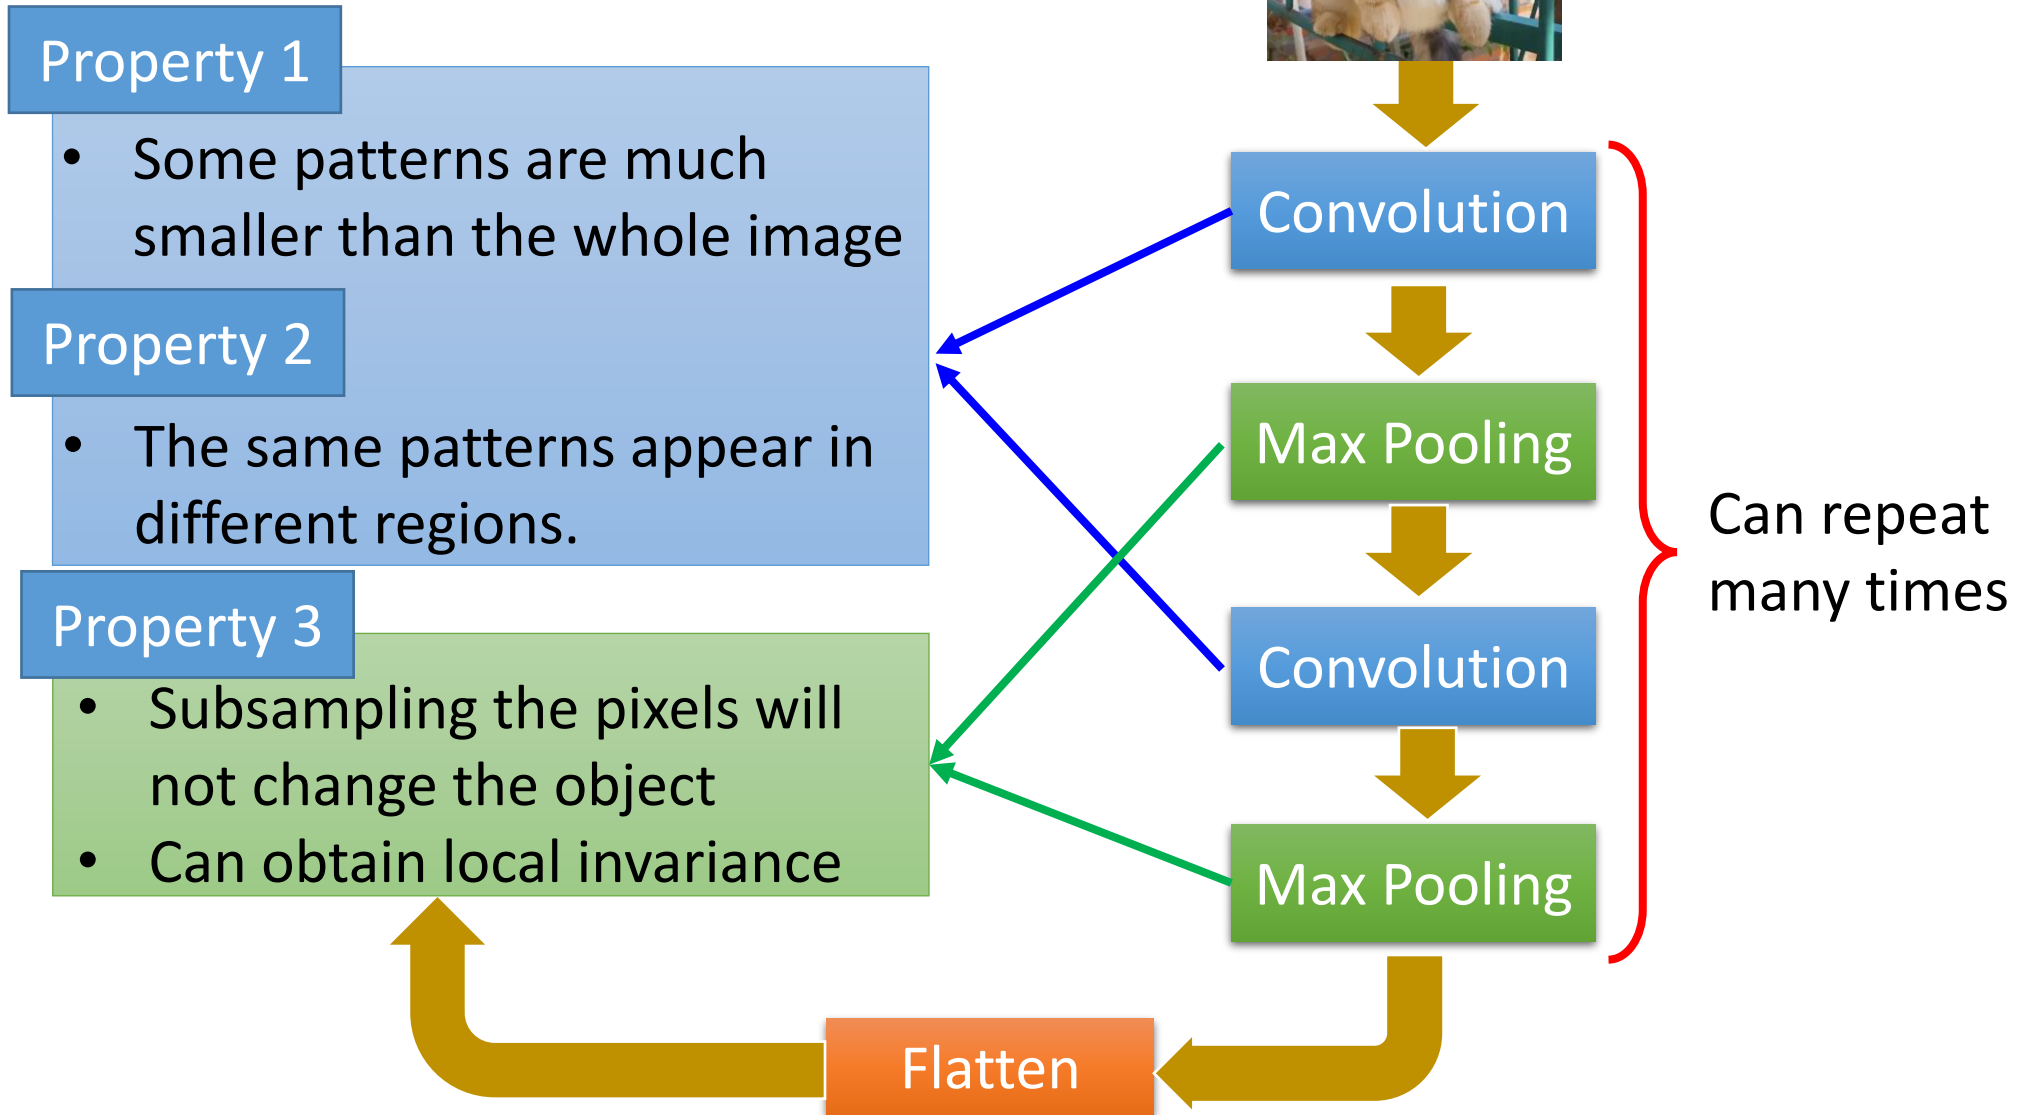
Max Pooling

Can repeat  
many times

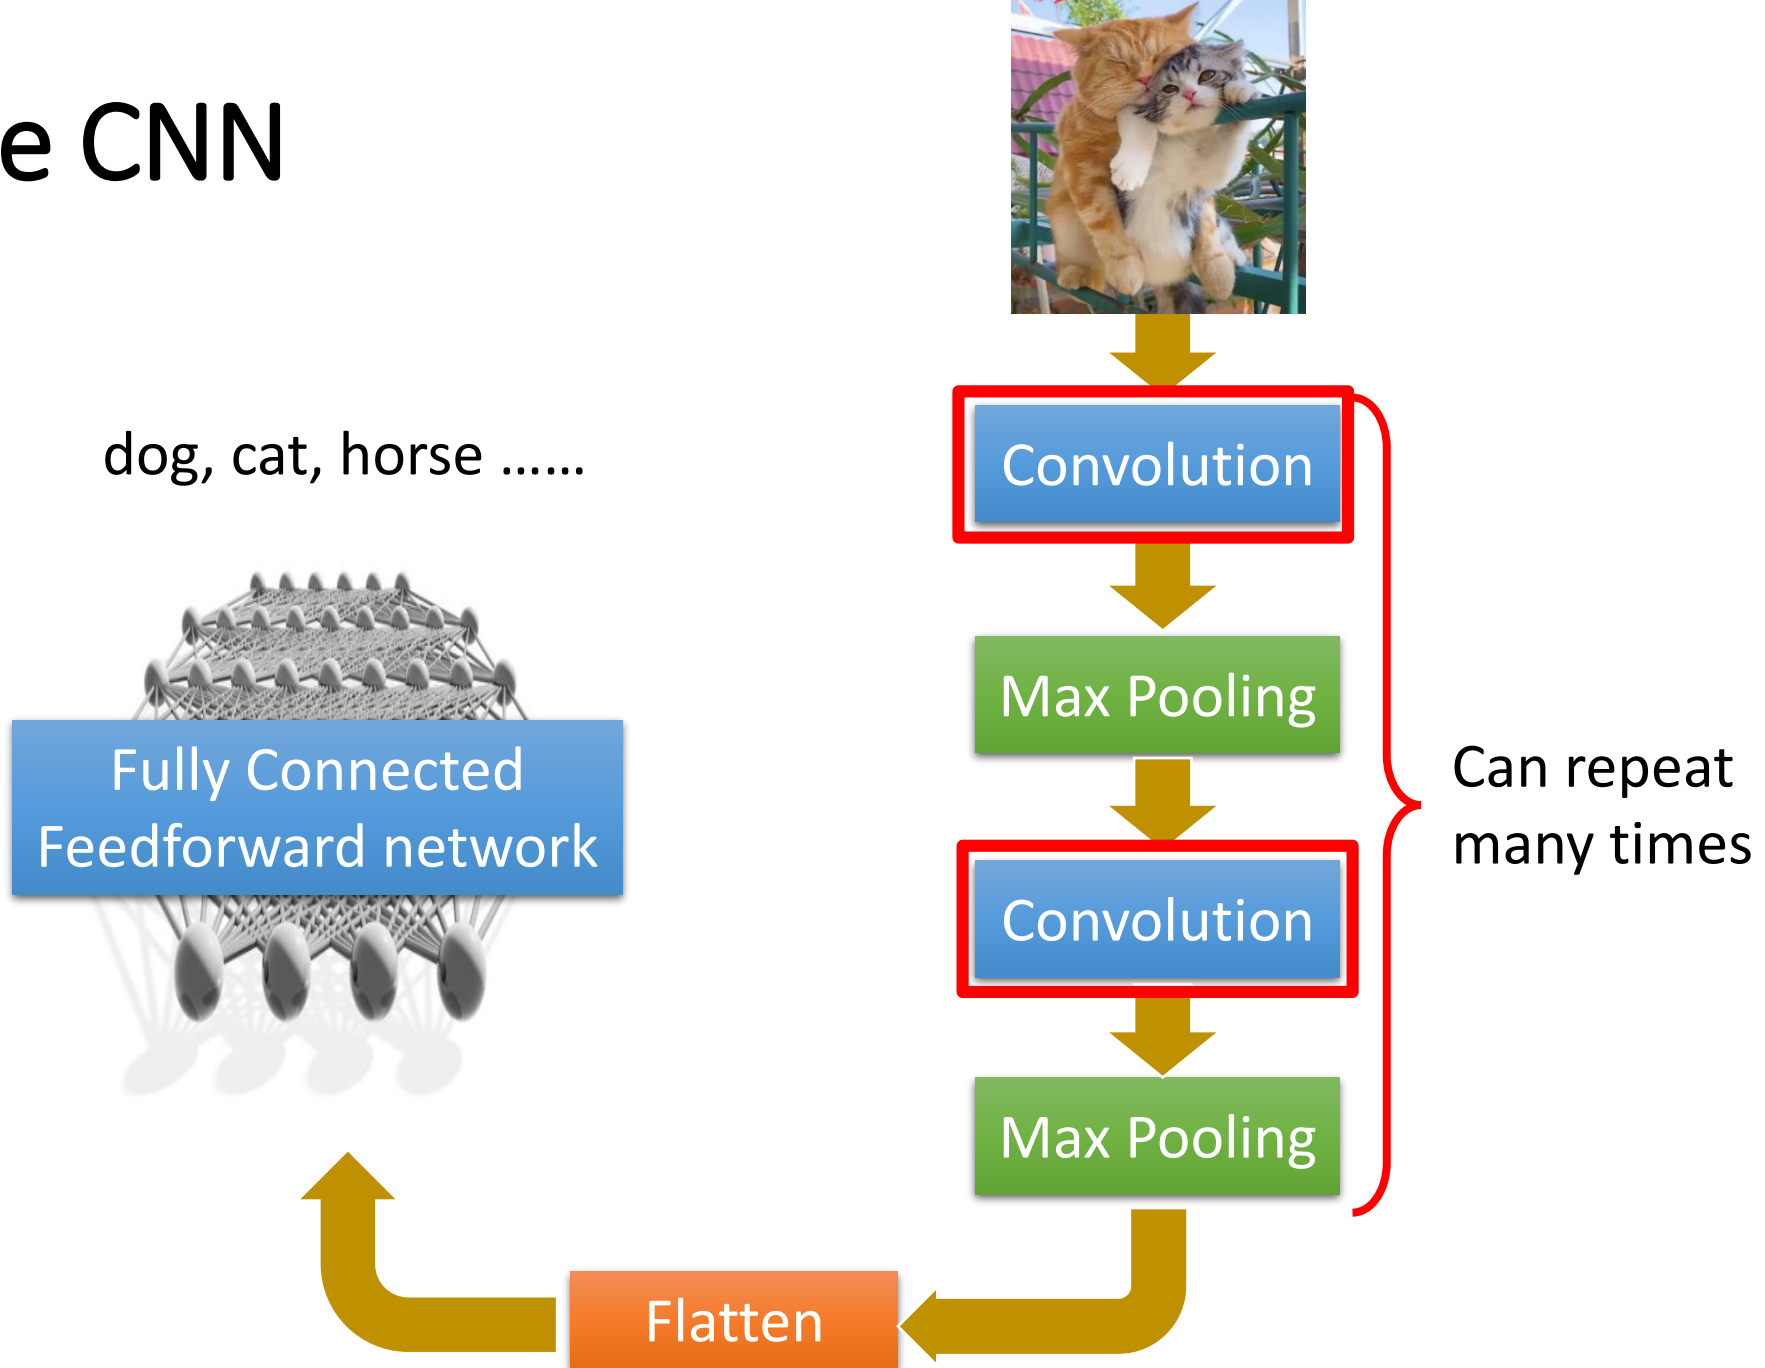
Flatten



# The whole CNN



# The whole CNN



# CNN – Convolution

Those are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 × 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1  
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2  
Matrix

⋮

Property 1

Each filter detects a small pattern (3 × 3).

# CNN – Convolution


stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 × 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



# CNN – Convolution


If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 × 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



We set stride=1 below

# CNN – Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 × 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Property 2

# CNN – Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 × 6 image

-1	1	-1
-1	1	-1
-1	1	-1

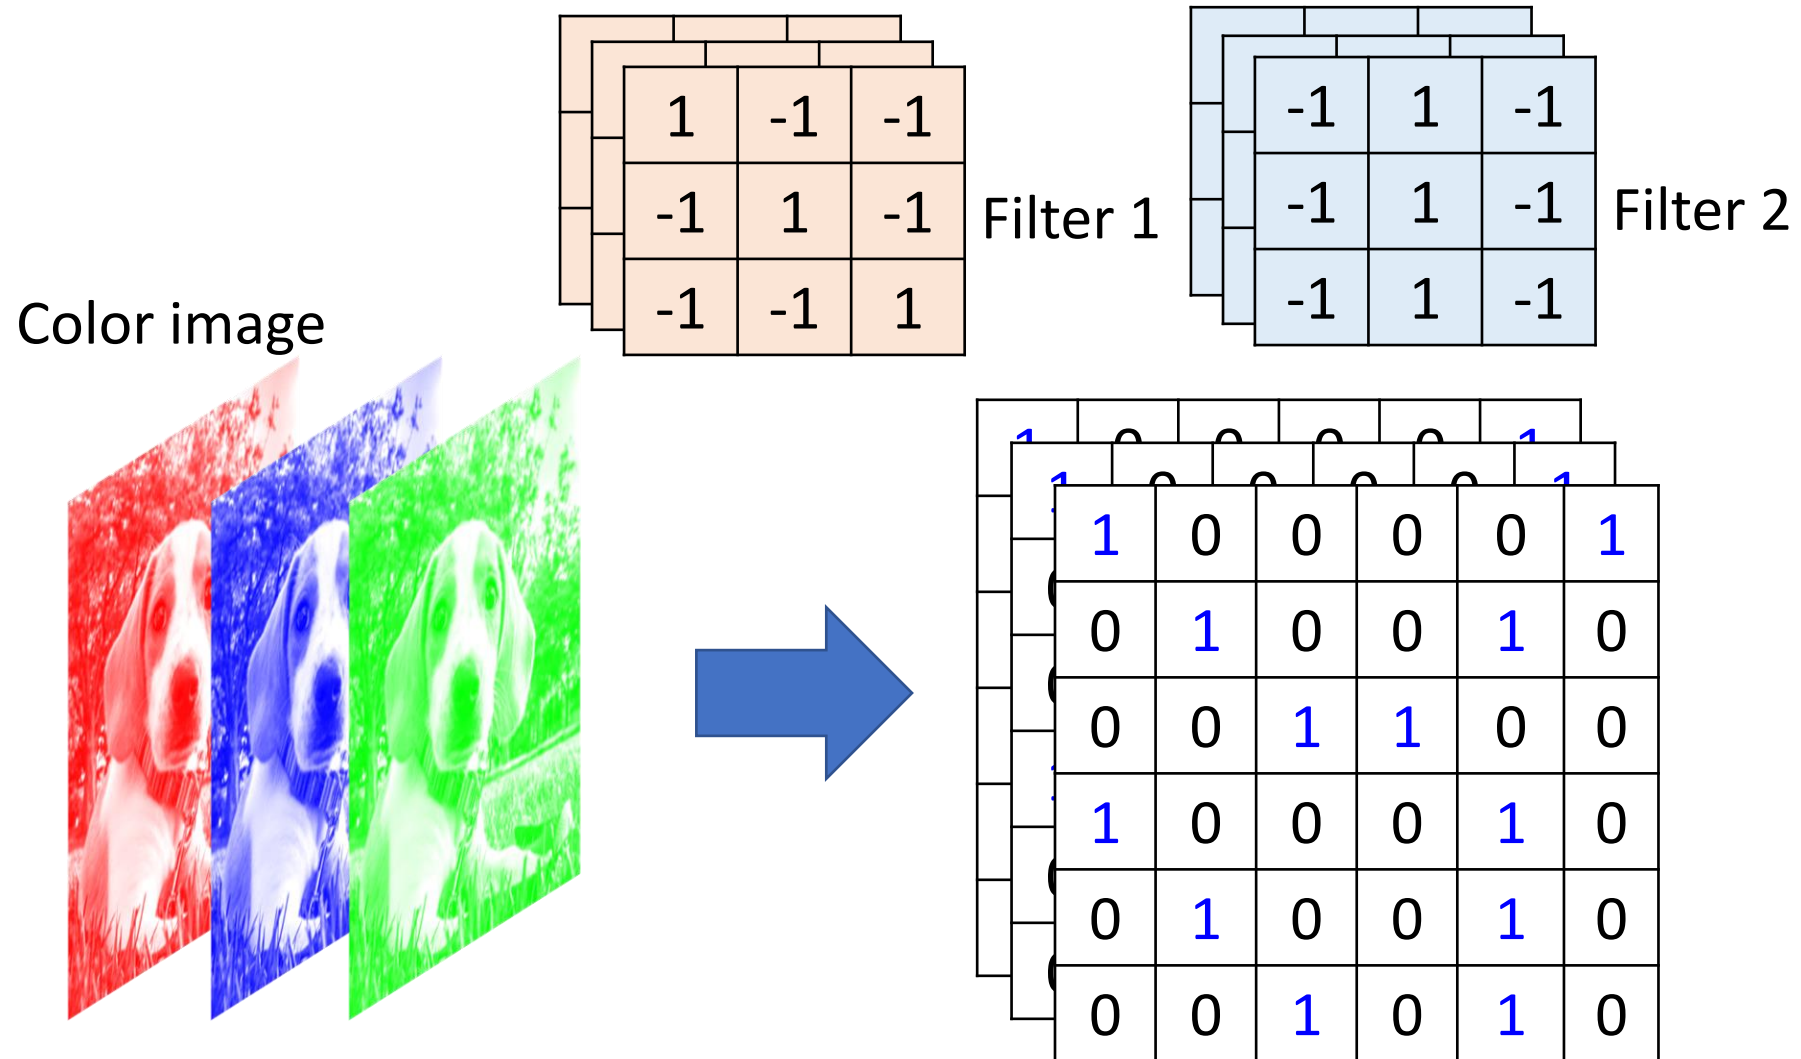
Filter 2

Do the same process for  
every filter




4 × 4 image

# CNN – Color Image







# Convolution v.s. Fully Connected




Fully Connected


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0






Less parameters!







Less parameters!

Even less parameters!



# The whole CNN

dog, cat, horse .....



Convolution

Max Pooling

Convolution

Max Pooling

Can repeat  
many times

Flatten

# CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

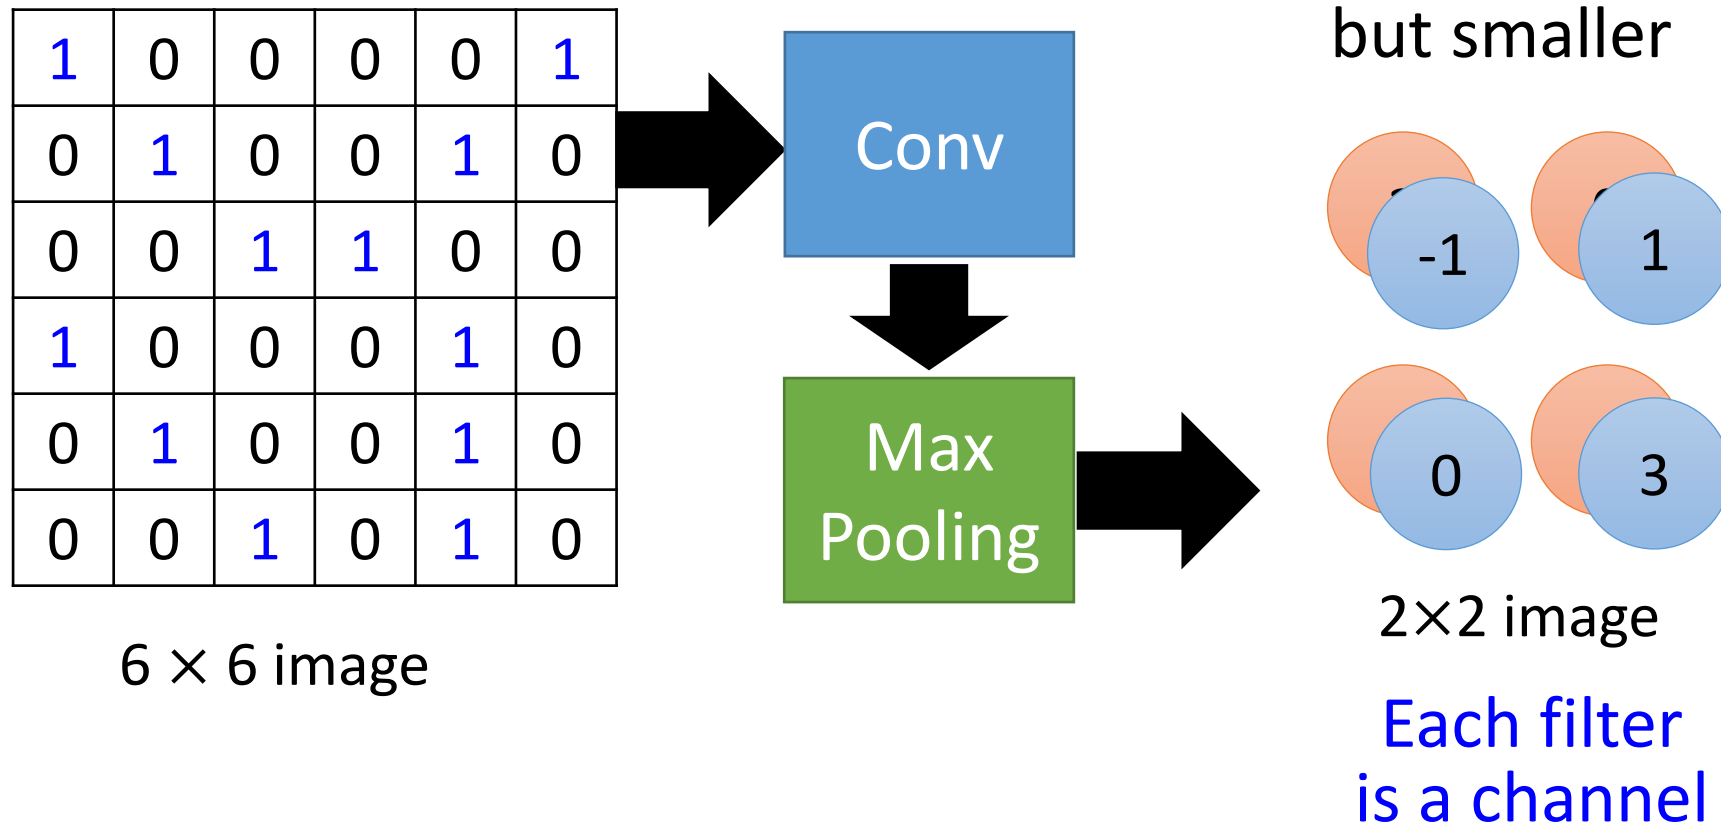
-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

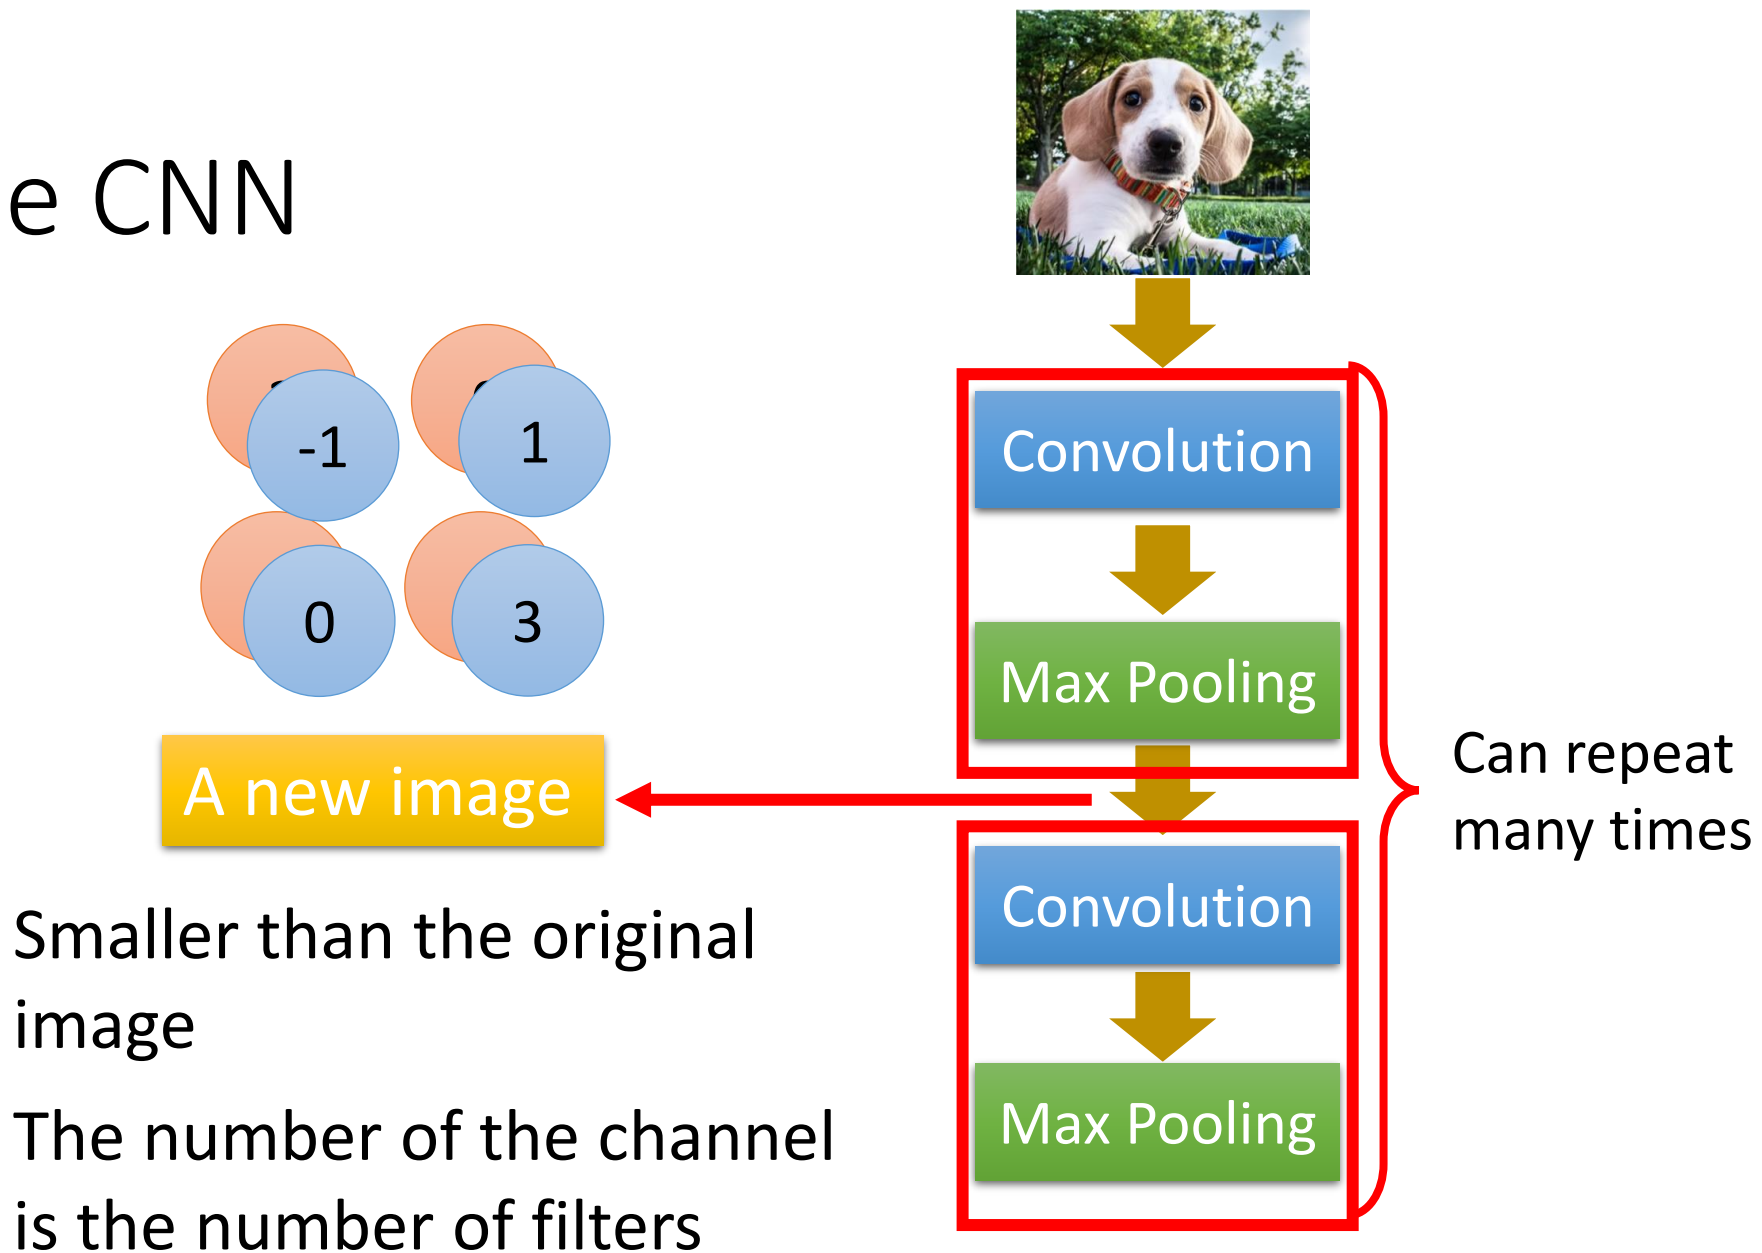
3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

# CNN – Max Pooling

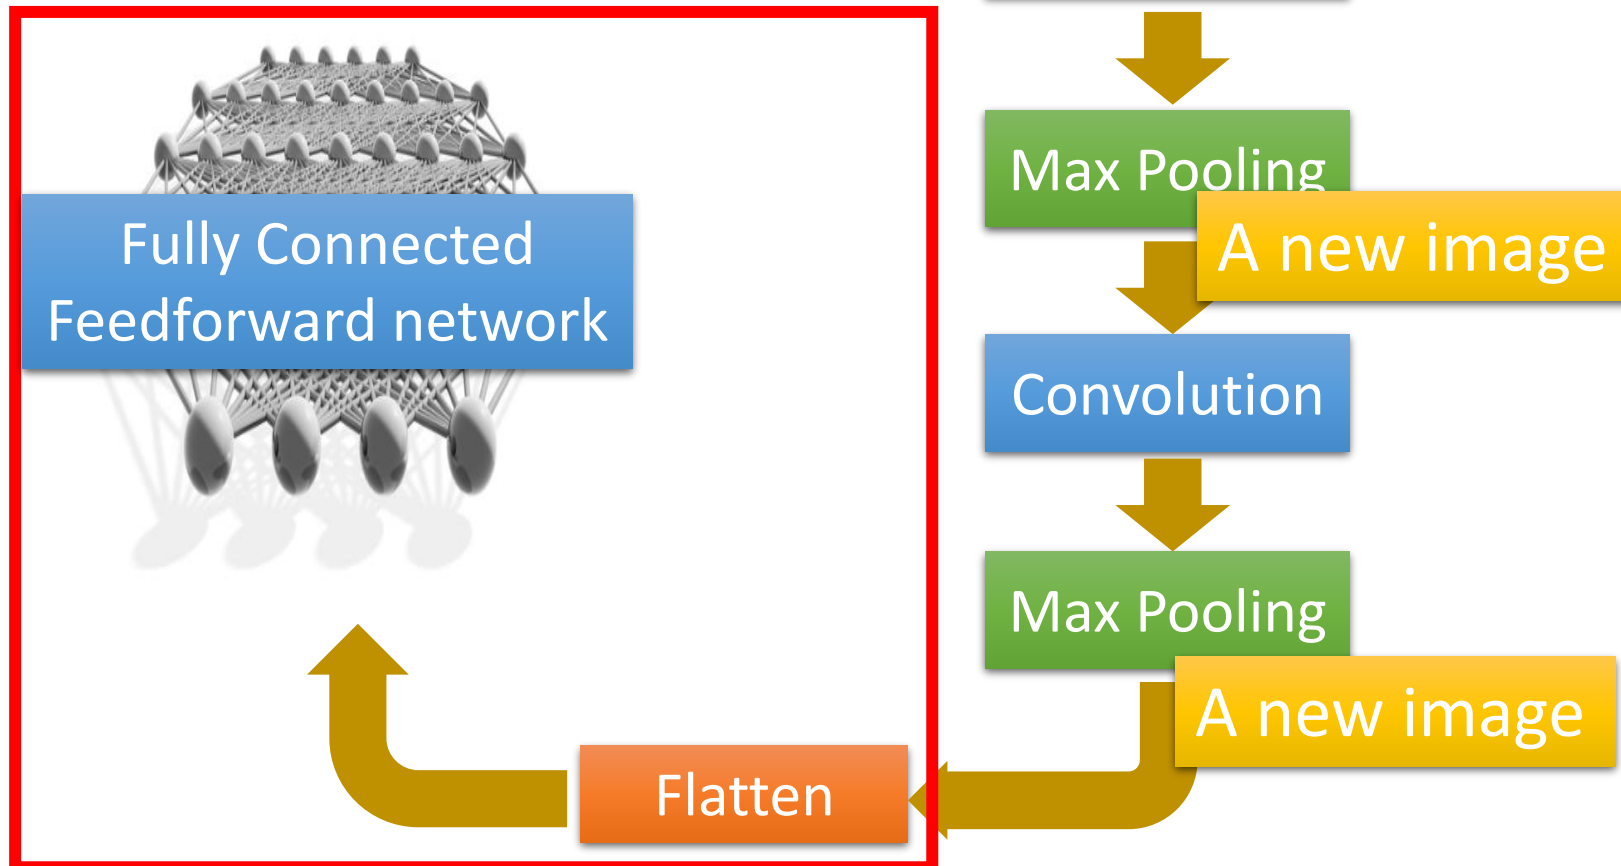


# The whole CNN



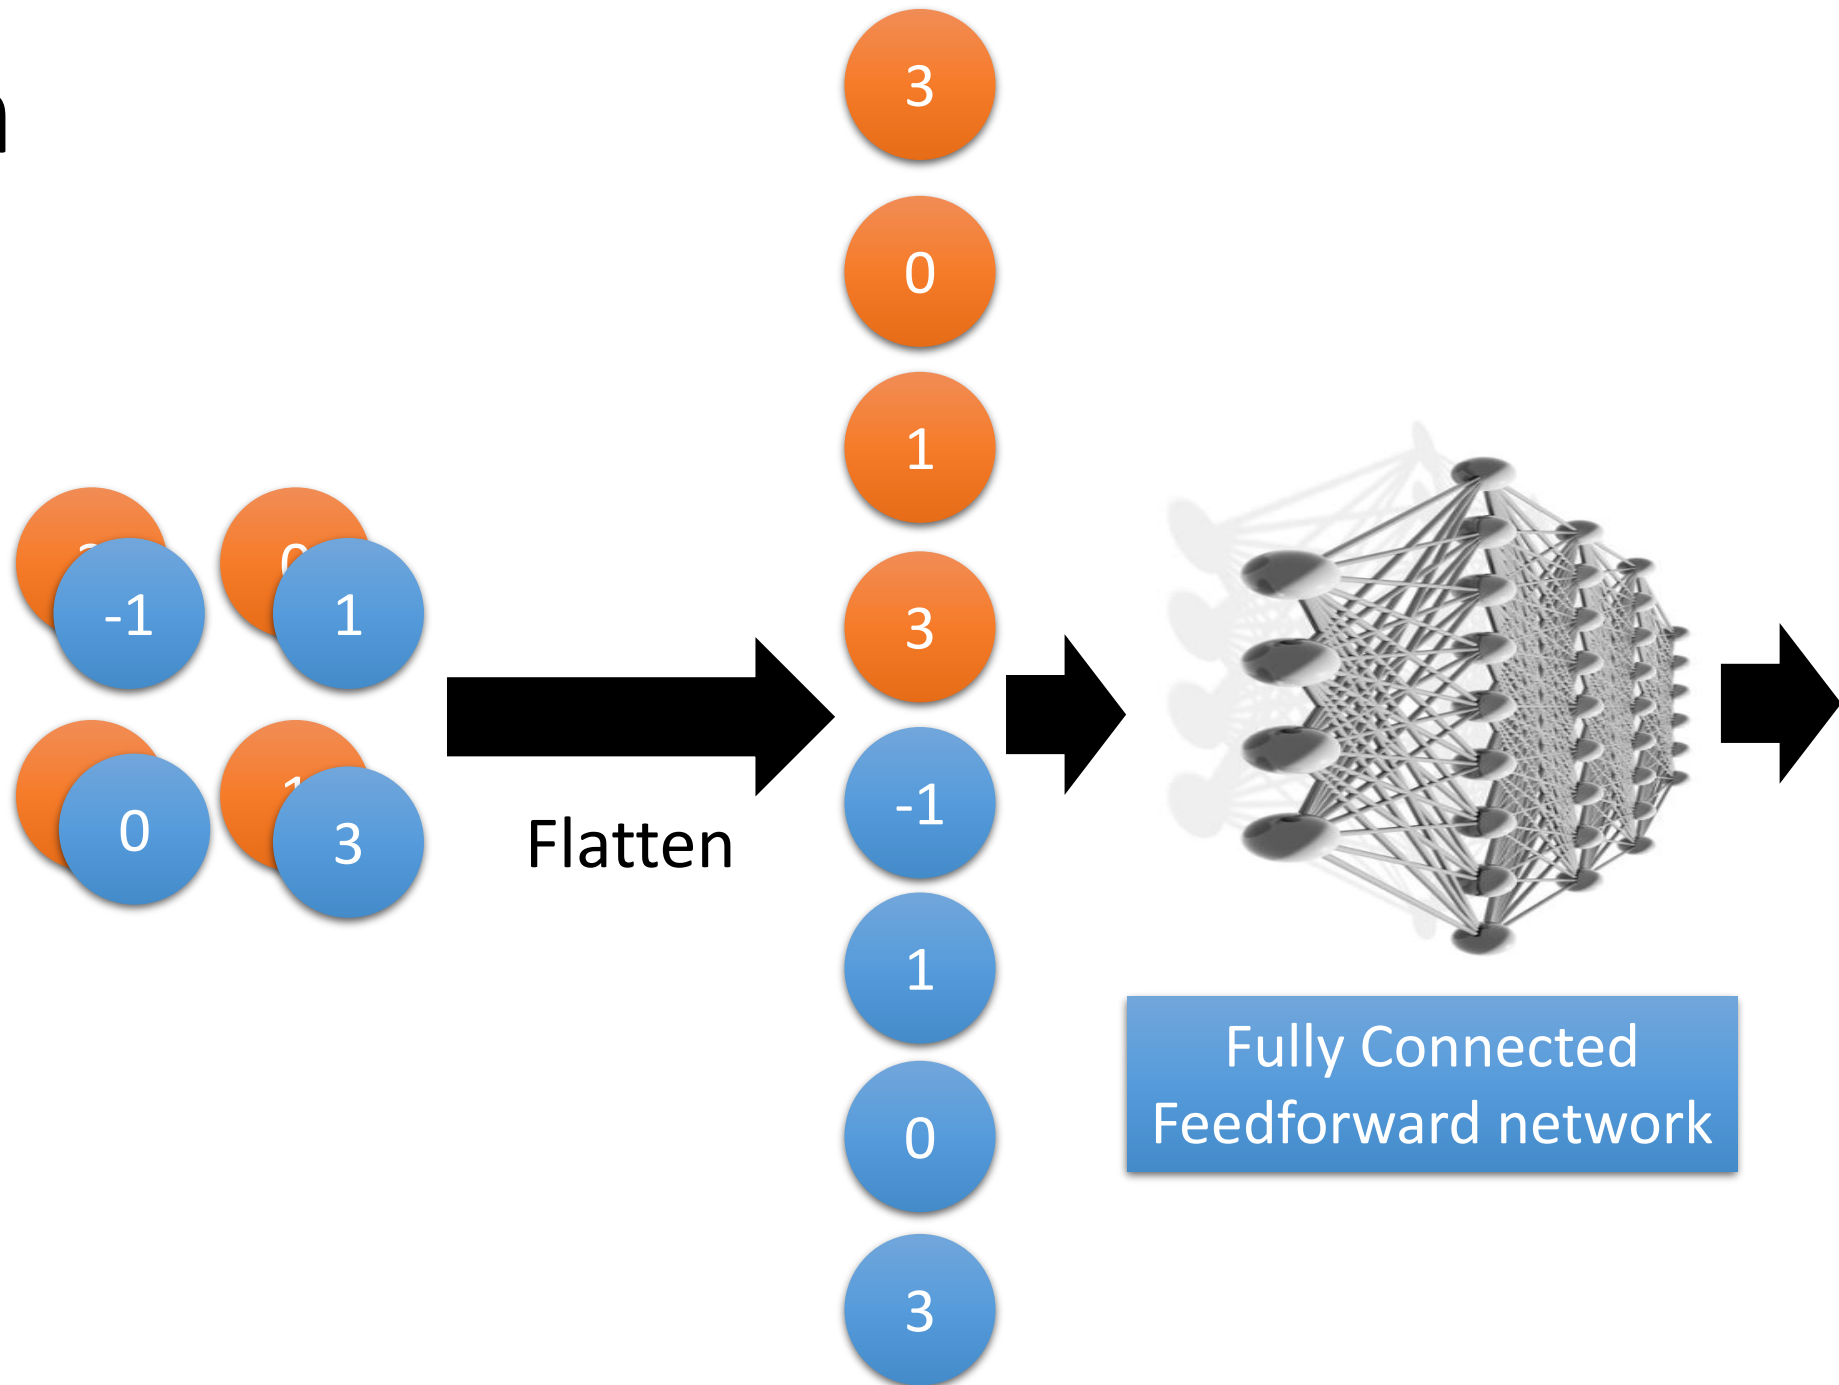
# The whole CNN

dog, cat, horse .....






# Flatten



# CNN-Summary

## ▣ Problems of fully connected neural networks


- Every output unit interacts with every input unit (pixel)
- The number of weights grows largely with the size of the input image
- Pixels in distance are less correlated



# CNN-Summary

## □ Locally connected neural networks


- Sparse connectivity: a hidden unit is only connected to a local patch (weights connected to the patch are called filter or kernel)
- Inspired by biological systems, where a cell is sensitive to a small subregion of the input space, called a receptive field. Many cells are tiled to cover the entire visual field.
- The design of such sparse connectivity is based on domain knowledge.



# CNN-Summary

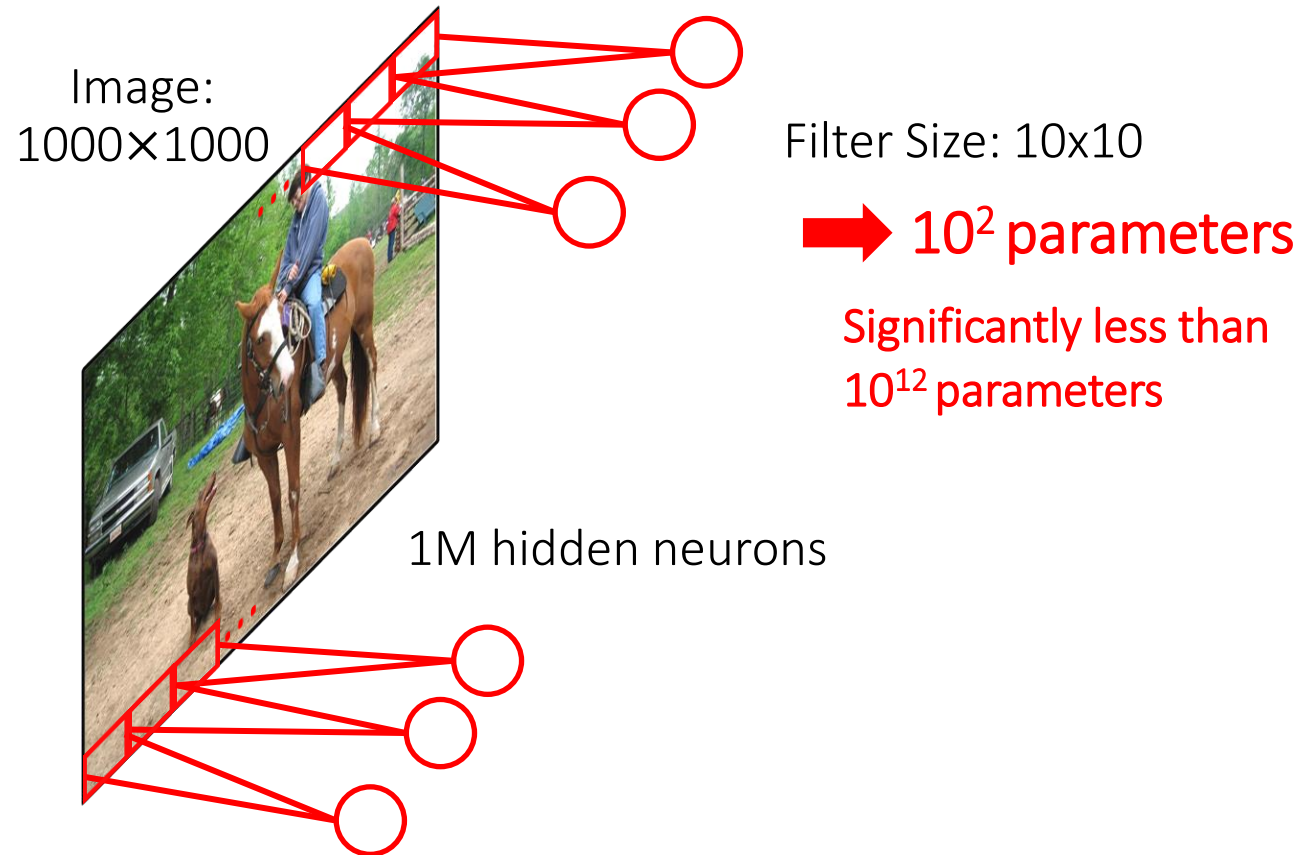
## ▣ Locally connected neural networks

- The learned filter is a spatially local pattern
- A hidden node at a higher layer has a larger receptive field in the input
- Stacking many such layers leads to “filters” which become increasingly “global”



# □ Share weights CNN-Summary


- Translation invariance: capture statistics in local patches and they are independent of locations. (Similar edges appear at different locations)
- Hidden nodes at different locations share the same weights. It greatly reduces the number of parameters to learn.
- We may only locally share weights or not share weights at top layers.



# CNN: Practices


## Convolution Layer

$32 \times 32 \times 3$  image




# Convolution Layer

$32 \times 32 \times 3$  image




$5 \times 5 \times 3$  filter



**Convolve** the filter with the image *i.e.* “slide over the image spatially, computing dot products”


# Convolution Layer

- $32 \times 32 \times 3$  image



Filters always extend the full depth of the input volume


$5 \times 5 \times 3$  filter



**Convolve** the filter with the image *i.e.* “slide over the image spatially, computing dot products”




# Convolution Layer



$$\mathbf{w}^T \mathbf{x} + b$$


# Convolution Layer



$7 \times 7$  input (spatially)  
assume  $3 \times 3$  filter


$\Rightarrow 5 \times 5$  output

# Convolution Layer




# Convolution Layer

consider a second, **green** filter




# Convolution Layer

For example, if we had 6 filters of size  $5 \times 5 \times 3$ , we'll get 6 separate activation maps:




We stack these up to get a “new image” of size  $28 \times 28 \times 6$ !


**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions.



**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions.




$32 \times 32$  input convolved repeatedly with  $3 \times 3$  filters shrinks volumes spatially!  
**Shrinking too fast is not good**, doesn't work well.





## A closer look at spatial dimensions:




7 × 7 input (spatially)  
assume 3 × 3 filter  
applied **with stride 2**

⇒ 3×3 output

$$\frac{\text{ImageSize} - \text{FilterSize}}{\text{Stride}} + 1$$

$$\frac{7-3}{2} + 1 = 3$$

A closer look at spatial dimensions:



$7 \times 7$  input (spatially)  
 assume  $3 \times 3$  filter  
 applied **with stride 3**?

**doesn't fit!**  
 cannot apply  $3 \times 3$  filter on  
 $7 \times 7$  input with stride 3.

# In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input  $7 \times 7$

$3 \times 3$  filter, applied with **stride 1**

**pad with 1 pixel border**  $\Rightarrow$  what is the output?

**$7 \times 7$  output!**

# In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input  $7 \times 7$

$3 \times 3$  filter, applied with **stride 1**

**pad with 1 pixel border**  $\Rightarrow$  what is the output?

**$7 \times 7$  output!**

in general, common to see CONV layers with stride 1, filters of size  $F \times F$ , and zero-padding with  $(F - 1)/2$ . (will preserve size spatially)


e.g.  $F = 3 \rightarrow$  zero pad with 1

$F = 5 \rightarrow$  zero pad with 2


$F = 7 \rightarrow$  zero pad with 3

# Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently



# MAX Pooling



# Lecture 4

## Convolutional Neural Networks

- CNN Basics
- Typical CNN Architectures

# Typical CNN Architectures

## Case Studies

- AlexNet
- VGG
- GoogLeNet
- ResNet

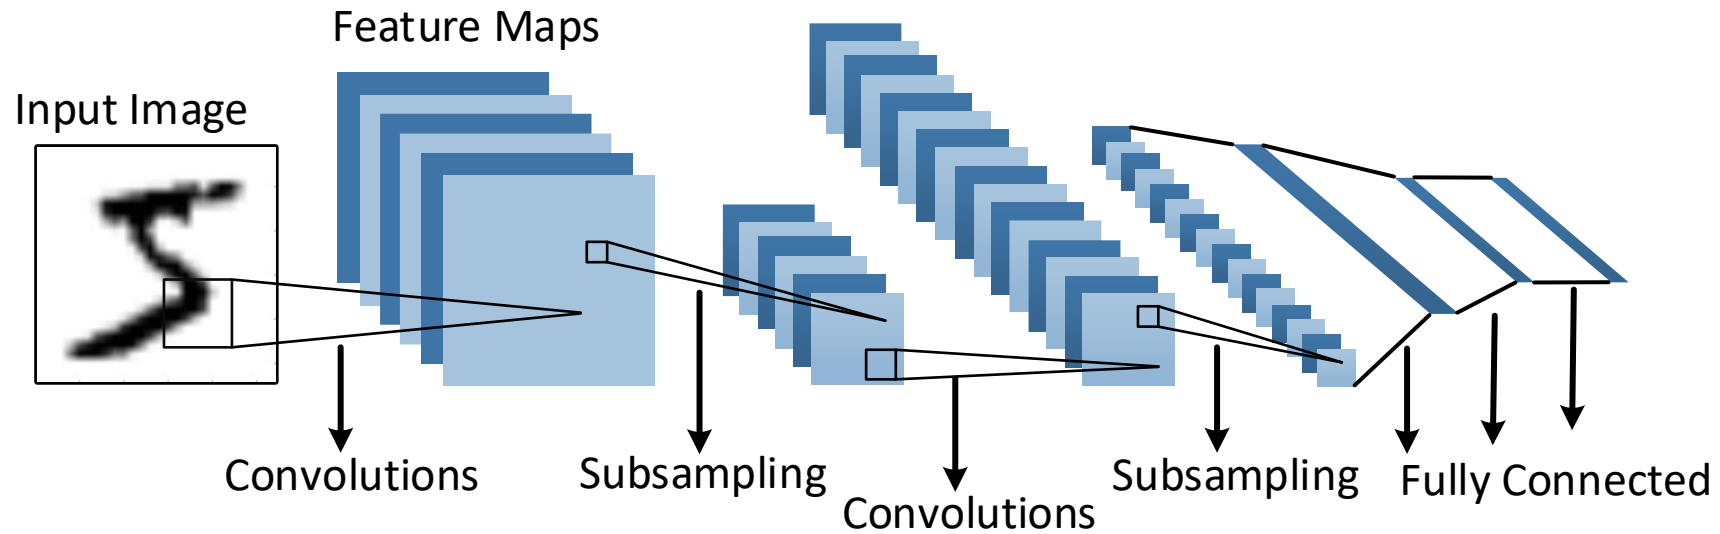
Briefly talk about ...(if time allows)

- NIN (Network in Network)
- Wide ResNet
- ResNeXT
- Stochastic Depth
- DenseNet
- SENet
- FractalNet
- SqueezeNet



# Review: LeNet

[LeCun *et al.* 1998]



Conv filters were  $5 \times 5$ , applied at stride 1

Subsampling (Pooling) layers were  $2 \times 2$  applied at stride 2

*i.e.* architecture is [CONV  $\rightarrow$  POOL  $\rightarrow$  CONV  $\rightarrow$  POOL  $\rightarrow$  FC  $\rightarrow$  FC]

# Case Study: AlexNet

[Krizhevsky *et al.* 2012]

Architecture:

CONV1

MAX POOL1

NORM1

CONV2

MAX POOL2

NORM2

CONV3

CONV4


CONV5

Max POOL3

FC6


FC7

FC8



# Case Study: AlexNet

[Krizhevsky *et al.* 2012]



Input: 227x227 x 3 images

**First layer (CONV1):** 96 11 x 11 filters applied at stride 4


Q: what is the output volume size?

$$\frac{\text{ImageSize} - \text{FilterSize}}{\text{Stride}} + 1$$

$$(227-11)/4+1 = 55$$

# Case Study: AlexNet

[Krizhevsky *et al.* 2012]




Input: 227×227×3 images

**First layer (CONV1):** 96 11×11 filters applied at stride 4  
Output volume [55×55×96]

Q: What is the total number of parameters in this layer?

# Case Study: AlexNet

[Krizhevsky *et al.* 2012]



Input: 227×227×3 images


**First layer (CONV1):** 96 11×11 filters applied at stride 4

Output volume **[55×55×96]**

Parameters:  $(11*11*3)*96 = \mathbf{35K}$

# Case Study: AlexNet

[Krizhevsky *et al.* 2012]



Input: 227×227×3 images


After CONV1: 55×55×96

**Second layer (POOL1):** 3×3 filters applied at stride 2

Q: what is the output volume size? Hint:  $(55-3)/2+1 = 27$

# Case Study: AlexNet

[Krizhevsky *et al.* 2012]



Input:  $227 \times 227 \times 3$  images

After CONV1:  $55 \times 55 \times 96$

**Second layer (POOL1):**  $3 \times 3$  filters applied at stride 2


Output volume:  $27 \times 27 \times 96$

Q: what is the number of parameters in this layer?

Parameters: 0!

# Case Study: AlexNet

[Krizhevsky *et al.* 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)

## Details:


- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- SGD Momentum 0.9
- Learning rate 0.01, reduced by 10 manually when val accuracy plateaus
- L2 weight decay 0.0005

7 CNN ensemble: 18.2% → 15.4%



# Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x48] x 2

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

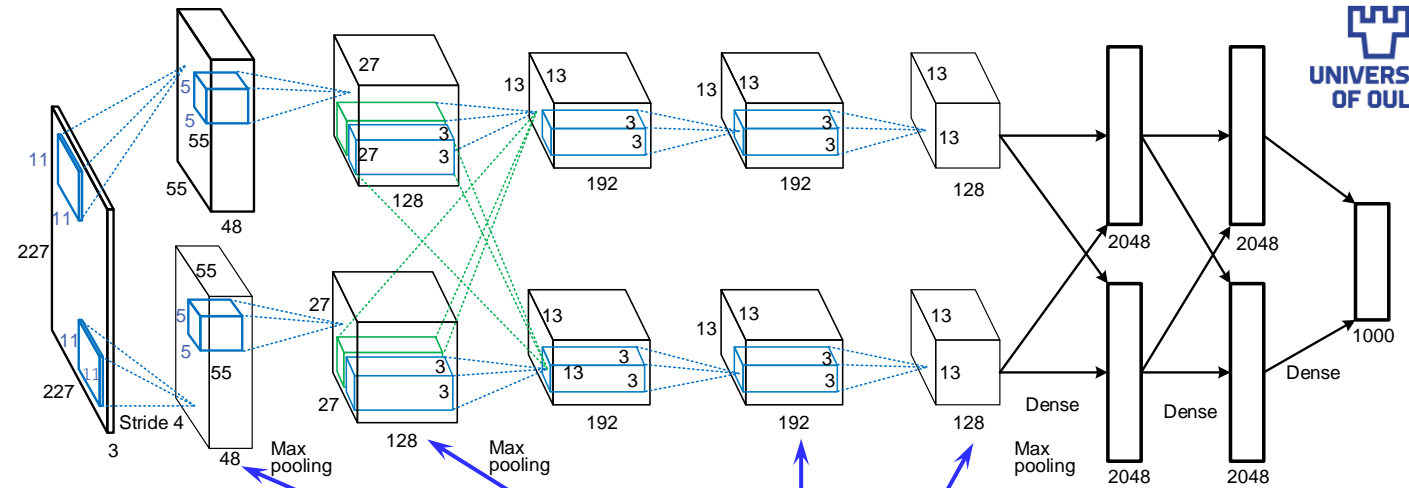
[1000] **FC8**: 1000 neurons (class scores)

Historical note: Trained on GTX 580  
GPU with only 3 GB of memory.

Network spread across 2 GPUs, half the  
neurons (feature maps) on each GPU.

# Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

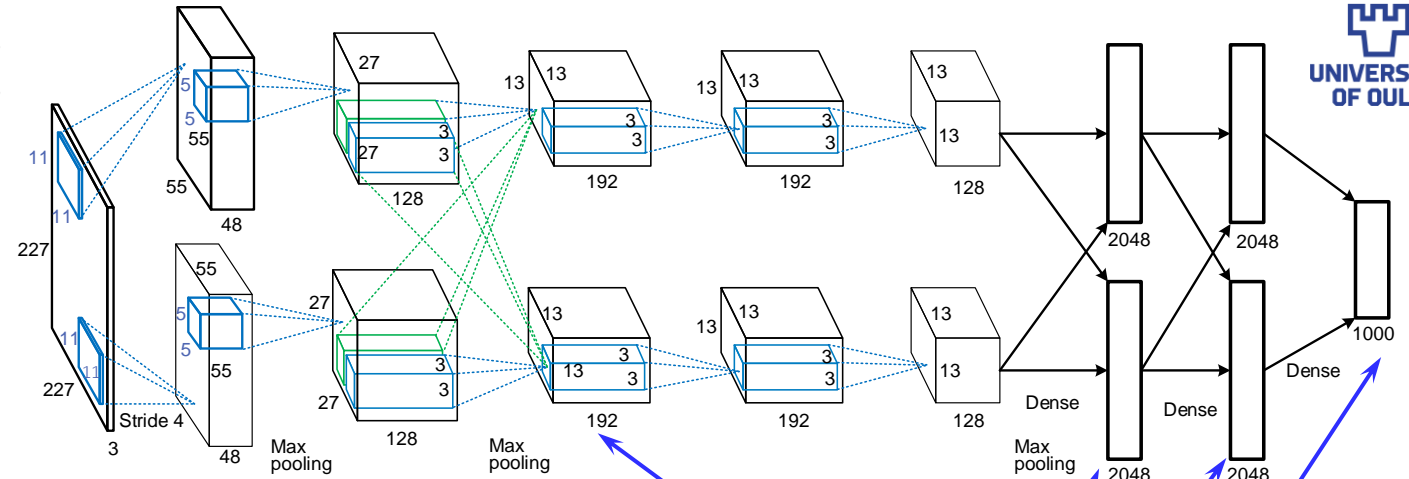
[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)

CONV1, CONV2, CONV4, CONV5:  
Connections only with feature maps  
on same GPU

# Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1


[6x6x256] **MAX POOL3**: 3x3 filters at stride 2


[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)


**CONV3, FC6, FC7, FC8:**  
Connections with all feature  
maps in preceding layer,  
communication across GPUs





# ZFNet


[Zeiler and Fergus, 2013]



AlexNet but: CONV1: change from (11x11 stride 4) to (7x7 stride 2)

CONV3,4,5: instead of 384, 384, 256 filters use 512, 1024, 512

ImageNet top 5 error: 16.4% → 11.7%



# Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Small filters, Deeper networks


8 layers (AlexNet)

→ 16–19 layers (VGG16Net)

Only  $3 \times 3$  CONV stride 1, pad 1 and  $2 \times 2$  MAX POOL stride 2

11.7% top 5 error in ILSVRC'13 (ZFNet)

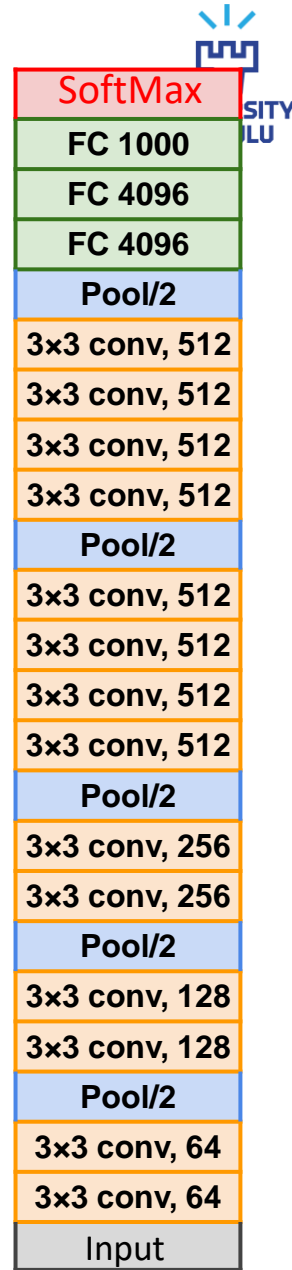
→ 7.3% top 5 error in ILSVRC'14



AlexNet



VGG16



VGG19




# Case Study: VGGNet

[Simonyan and Zisserman, 2014]


Q: Why use smaller filters? ( $3 \times 3$  conv)

Stack of three  $3 \times 3$  conv (stride 1) layers has same **effective receptive field** as one  $7 \times 7$  conv layer


Q: What is the effective receptive field of three  $3 \times 3$  conv (stride 1) layers?



AlexNet



VGG16





VGG19


# Case Study: VGGNet

[Simonyan and Zisserman, 2014]


Q: What is the effective receptive field of three  $3 \times 3$  conv (stride 1) layers?



AlexNet



VGG16





VGG19


# Case Study: VGGNet

[Simonyan and Zisserman, 2014]


Q: What is the effective receptive field of three  $3 \times 3$  conv (stride 1) layers?



AlexNet



VGG16





VGG19


# Case Study: VGGNet

[Simonyan and Zisserman, 2014]


Q: What is the effective receptive field of three  $3 \times 3$  conv (stride 1) layers?



AlexNet



VGG16





VGG19


# Case Study: VGGNet

[Simonyan and Zisserman, 2014]


Q: What is the effective receptive field of three  $3 \times 3$  conv (stride 1) layers?



AlexNet



VGG16



VGG19

# Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Q: Why use smaller filters? ( $3 \times 3$  conv)

Stack of three  $3 \times 3$  conv (stride 1) layers has same **effective receptive field** as one  $7 \times 7$  conv layer


Q: What is the effective receptive field of three  $3 \times 3$  conv (stride 1) layers?

$[7 \times 7]$


But deeper, more nonlinearities

And fewer parameters:


$3 \cdot (3^2 C^2)$  vs.  $7^2 C^2$  for  $C$  channels per layer



AlexNet



VGG16



VGG19

INPUT: [224x224x3] **memory:**  $224*224*3=150\text{K}$  **params:** 0

CONV3-64: [224x224x64] **memory:**  $224*224*64=3.2\text{M}$  **params:**  $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] **memory:**  $224*224*64=3.2\text{M}$  **params:**  $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] **memory:**  $112*112*64=800\text{K}$  **params:** 0

CONV3-128: [112x112x128] **memory:**  $112*112*128=1.6\text{M}$  **params:**  $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] **memory:**  $112*112*128=1.6\text{M}$  **params:**  $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] **memory:**  $56*56*128=400\text{K}$  **params:** 0

CONV3-256: [56x56x256] **memory:**  $56*56*256=800\text{K}$  **params:**  $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] **memory:**  $56*56*256=800\text{K}$  **params:**  $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] **memory:**  $56*56*256=800\text{K}$  **params:**  $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] **memory:**  $28*28*256=200\text{K}$  **params:** 0

CONV3-512: [28x28x512] **memory:**  $28*28*512=400\text{K}$  **params:**  $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] **memory:**  $28*28*512=400\text{K}$  **params:**  $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] **memory:**  $28*28*512=400\text{K}$  **params:**  $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] **memory:**  $14*14*512=100\text{K}$  **params:** 0

CONV3-512: [14x14x512] **memory:**  $14*14*512=100\text{K}$  **params:**  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] **memory:**  $14*14*512=100\text{K}$  **params:**  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] **memory:**  $14*14*512=100\text{K}$  **params:**  $(3*3*512)*512 = 2,359,296$


POOL2: [7x7x512] **memory:**  $7*7*512=25\text{K}$  **params:** 0

FC: [1x1x4096] **memory:** 4096 **params:**  $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] **memory:** 4096 **params:**  $4096*4096 = 16,777,216$

FC: [1x1x1000] **memory:** 1000 **params:**  $4096*1000 = 4,096,000$

fc8  
fc7  
fc6  
conv5-3  
conv5-2  
conv5-1  
conv4-3  
conv4-2  
conv4-1  
conv3-2  
conv3-1  
conv2-2  
conv2-1  
conv1-2  
conv1-1



VGG16

**TOTAL memory:**  $24\text{M} * 4 \text{ bytes} \sim 96\text{MB}$  / image (only forward! about\*2 for bwd)

**TOTAL params:** 138M parameters



INPUT: [224x224x3] **memory: 224\*224\*3=150K** params: 0

CONV3-64: [224x224x64] **memory: 224\*224\*64=3.2M** params:  $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] **memory: 224\*224\*64=3.2M** params:  $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] **memory: 112\*112\*64=800K** params: 0

CONV3-128: [112x112x128] **memory: 112\*112\*128=1.6M** params:  $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] **memory: 112\*112\*128=1.6M** params:  $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] **memory: 56\*56\*128=400K** params: 0

CONV3-256: [56x56x256] **memory: 56\*56\*256=800K** params:  $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] **memory: 56\*56\*256=800K** params:  $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] **memory: 56\*56\*256=800K** params:  $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] **memory: 28\*28\*256=200K** params: 0

CONV3-512: [28x28x512] **memory: 28\*28\*512=400K** params:  $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] **memory: 28\*28\*512=400K** params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] **memory: 28\*28\*512=400K** params:  $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] **memory: 14\*14\*512=100K** params: 0

CONV3-512: [14x14x512] **memory: 14\*14\*512=100K** params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] **memory: 14\*14\*512=100K** params:  $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] **memory: 14\*14\*512=100K** params:  $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] **memory: 7\*7\*512=25K** params: 0

FC: [1x1x4096] **memory: 4096** params:  **$7*7*512*4096 = 102,760,448$**

FC: [1x1x4096] **memory: 4096** params:  **$4096*4096 = 16,777,216$**

FC: [1x1x1000] **memory: 1000** params:  $4096*1000 = 4,096,000$

Most memory is  
in early CONV



SoftMax
FC 1000
FC 4096
FC 4096
Pool/2
3x3 conv, 512
3x3 conv, 512
3x3 conv, 512
Pool/2
3x3 conv, 512
3x3 conv, 512
3x3 conv, 512
Pool/2
3x3 conv, 256
3x3 conv, 256
Pool/2
3x3 conv, 128
3x3 conv, 128
Pool/2
3x3 conv, 64
3x3 conv, 64
Input

Most params  
are in late FC

**TOTAL memory: 24M \* 4 bytes ~ = 96MB / image** (only forward! about\*2 for bwd)

**TOTAL params: 138M parameters**

VGG16




# Case Study: VGGNet

[Simonyan and Zisserman, 2014]

## Summary


- ILSVRC'14 2nd in classification, 1st in localization
- Similar training procedure as AlexNet
- No Local Response Normalization (LRN)
- Use VGG16 or VGG19 (VGG19 only slightly better, more memory)
- Use ensembles for best results
- FC7 features generalize well to other tasks




AlexNet



VGG16



VGG19

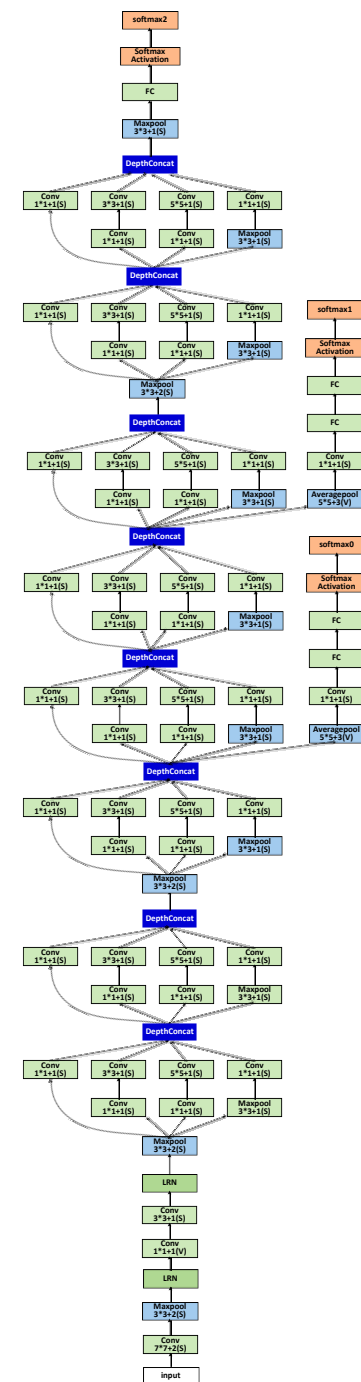


# Case Study: GoogLeNet

[Szegedy et al., 2014]

## Deeper networks and computational efficiency

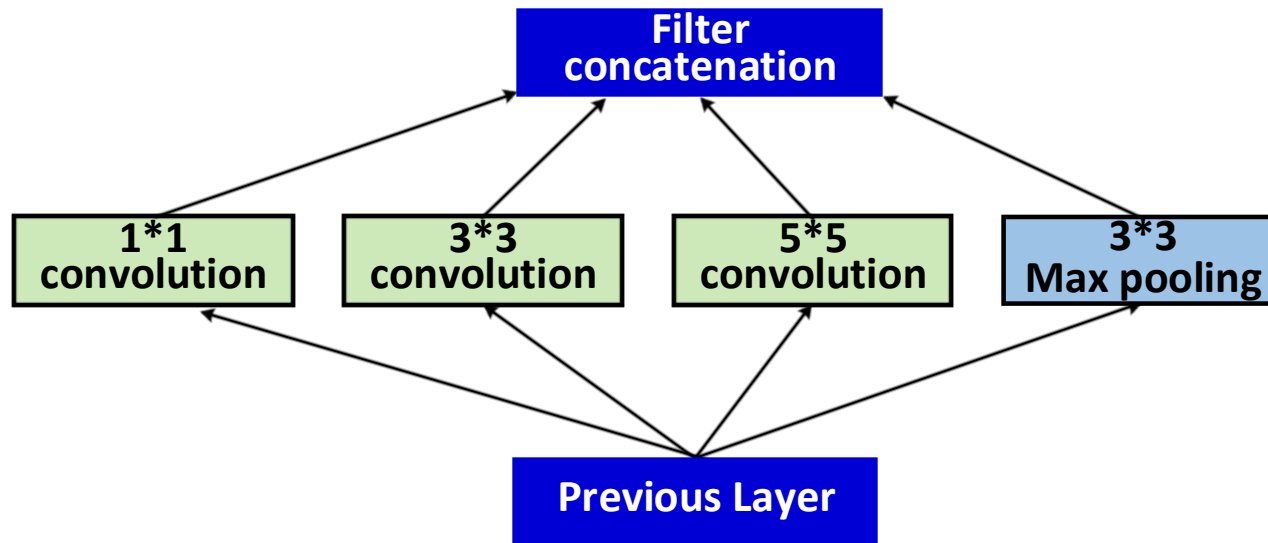
- 22 layers
- Efficient “Inception” module
- No FC layers
- Only 5 million parameters!
  - 12× less than AlexNet
- ILSVRC’14 classification winner (6.7% top 5 error)





# Case Study: GoogLeNet

[Szegedy et al., 2014]



Naive Inception module

Apply parallel filter operations on the input from previous layer:

- Multiple receptive field sizes for convolution ( $1 \times 1$ ,  $3 \times 3$ ,  $5 \times 5$ )
- Pooling operation ( $3 \times 3$ )

Concatenate all filter outputs together depthwise.

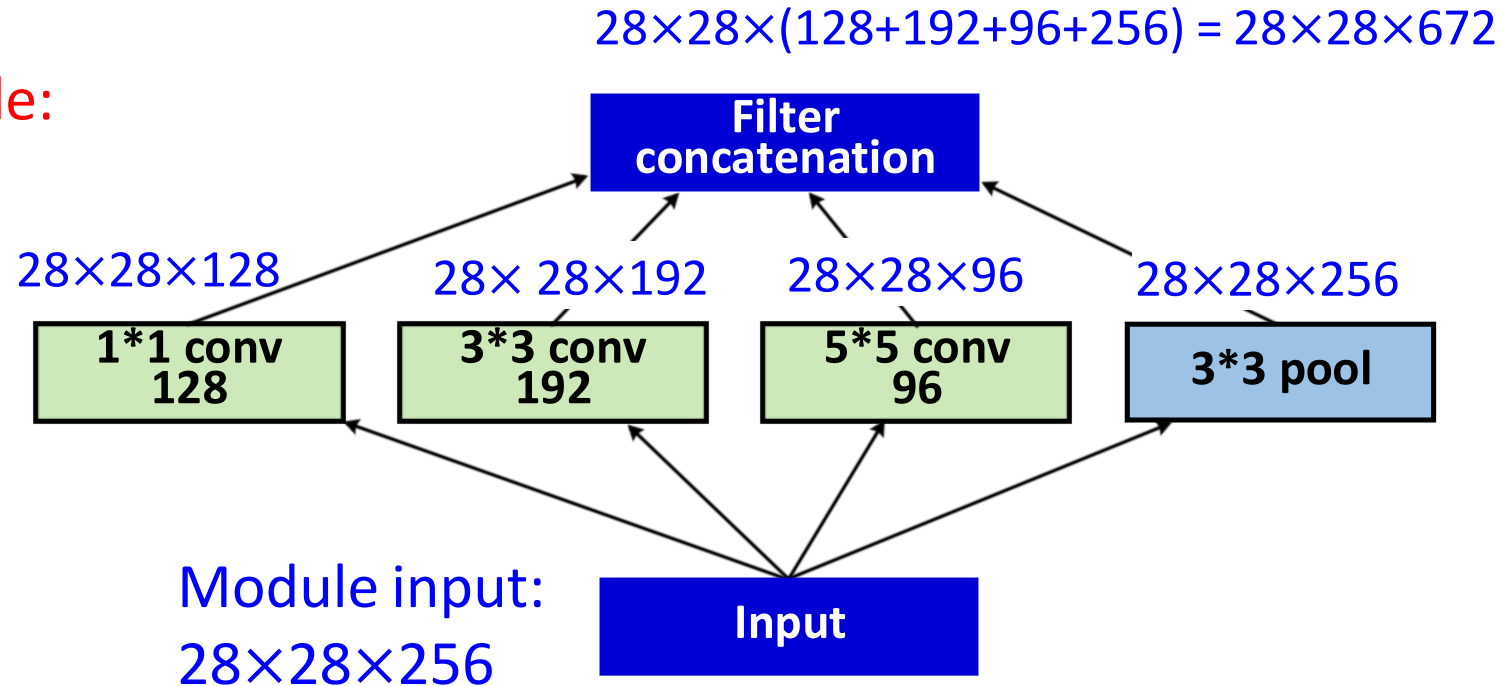
Q: What is the problem with this?  
[Hint: Computational complexity]

# Case Study: GoogLeNet

[Szegedy et al., 2014]

Q: What is the problem with this?  
[Hint: Computational complexity]

Example:



Naive Inception module

Q1: What is the output size of the  $1 \times 1$  conv, with 128 filters?

Q2: What are the output sizes of all different filter operations?

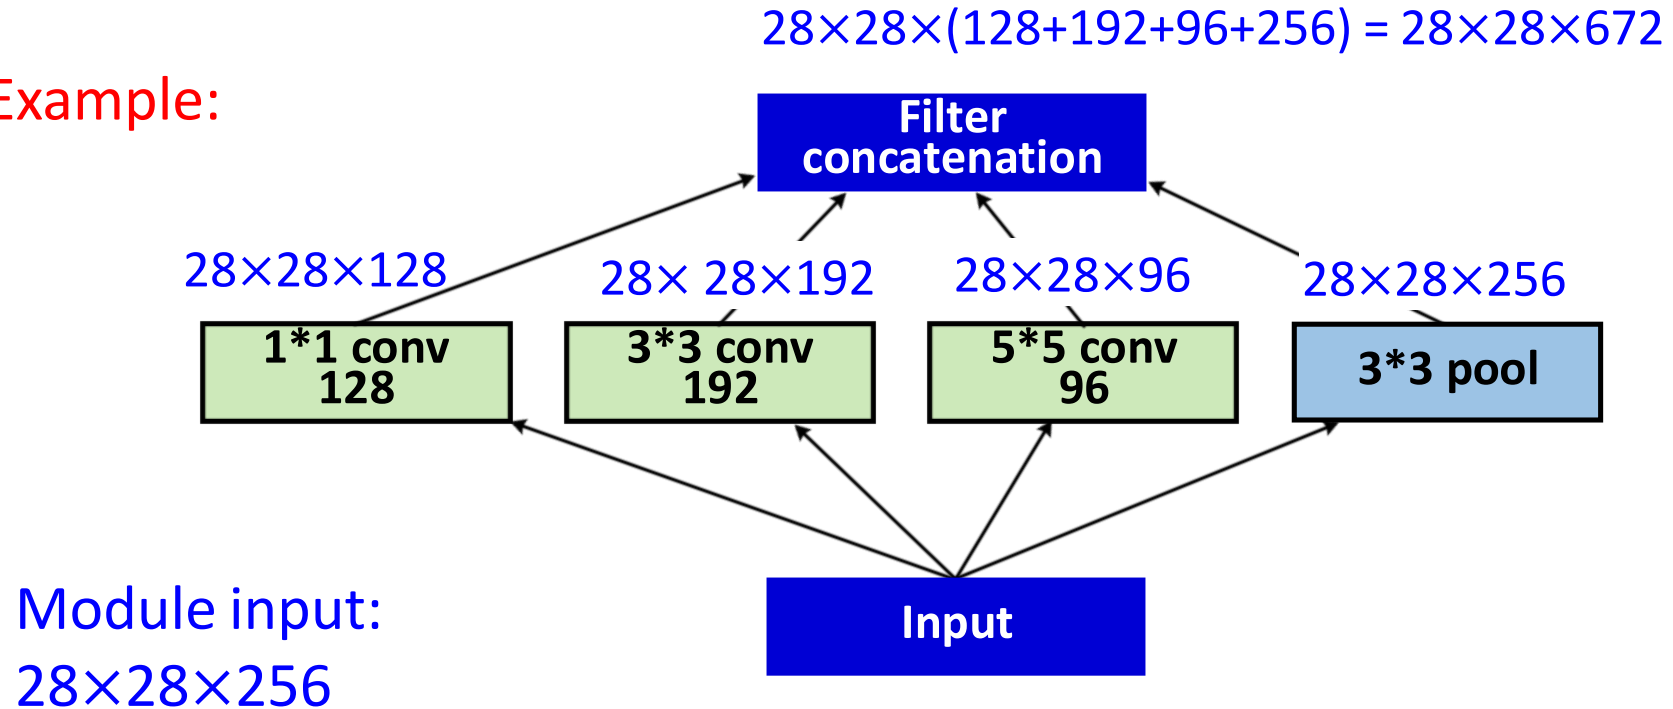
Q3: What is output size after filter concatenation?

# Case Study: GoogLeNet

Q: What is the problem with this?  
[Hint: Computational complexity]

[Szegedy et al., 2014]

Example:



## Naive Inception module

Very expensive compute

**Conv Ops:**

$[1 \times 1 \text{ conv, } 128] \ 28 \times 28 \times 128 \times 1 \times 1 \times 256$

$[3 \times 3 \text{ conv, } 192] \ 28 \times 28 \times 192 \times 3 \times 3 \times 256$

$[5 \times 5 \text{ conv, } 96] \ 28 \times 28 \times 96 \times 5 \times 5 \times 256$

**Total: 854M ops**

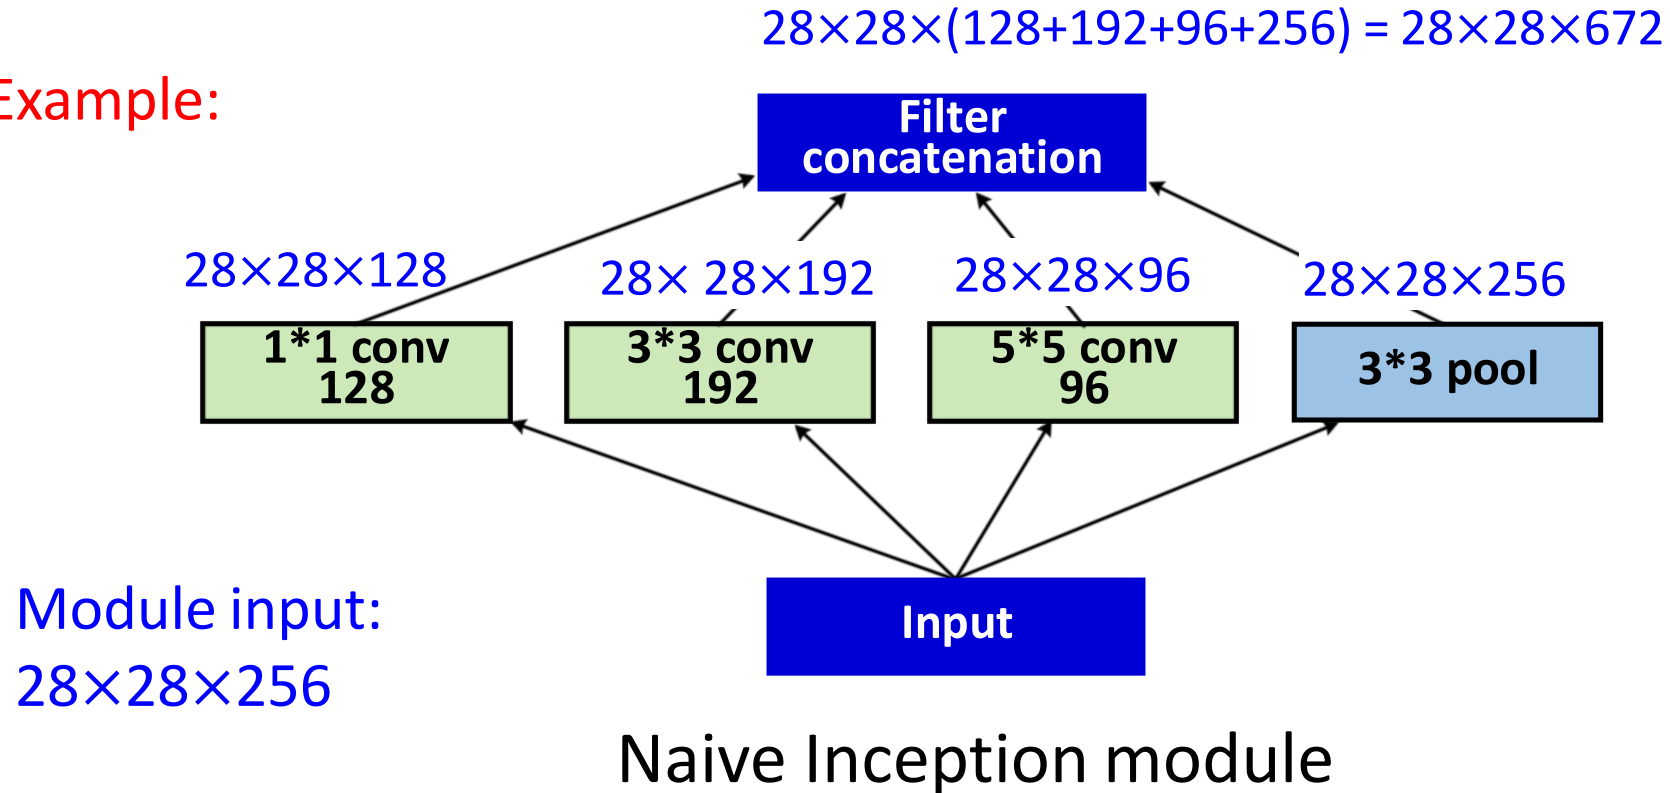
Pooling layer also preserves feature depth, which means total depth after concatenation can only grow at every layer!

# Case Study: GoogLeNet

Q: What is the problem with this?  
[Hint: Computational complexity]

[Szegedy et al., 2014]

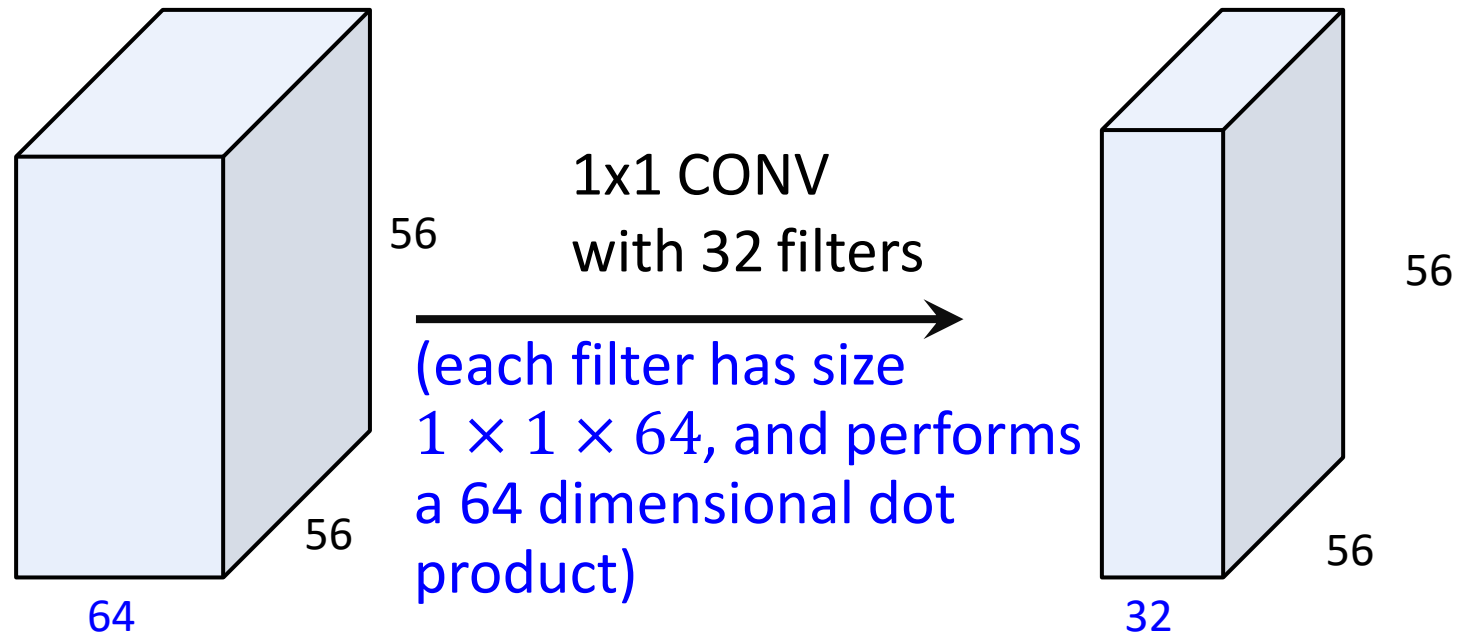
Example:



Solution: “bottleneck” layers that use  $1 \times 1$  convolutions to reduce feature depth.



## Reminder: 1x1 convolutions

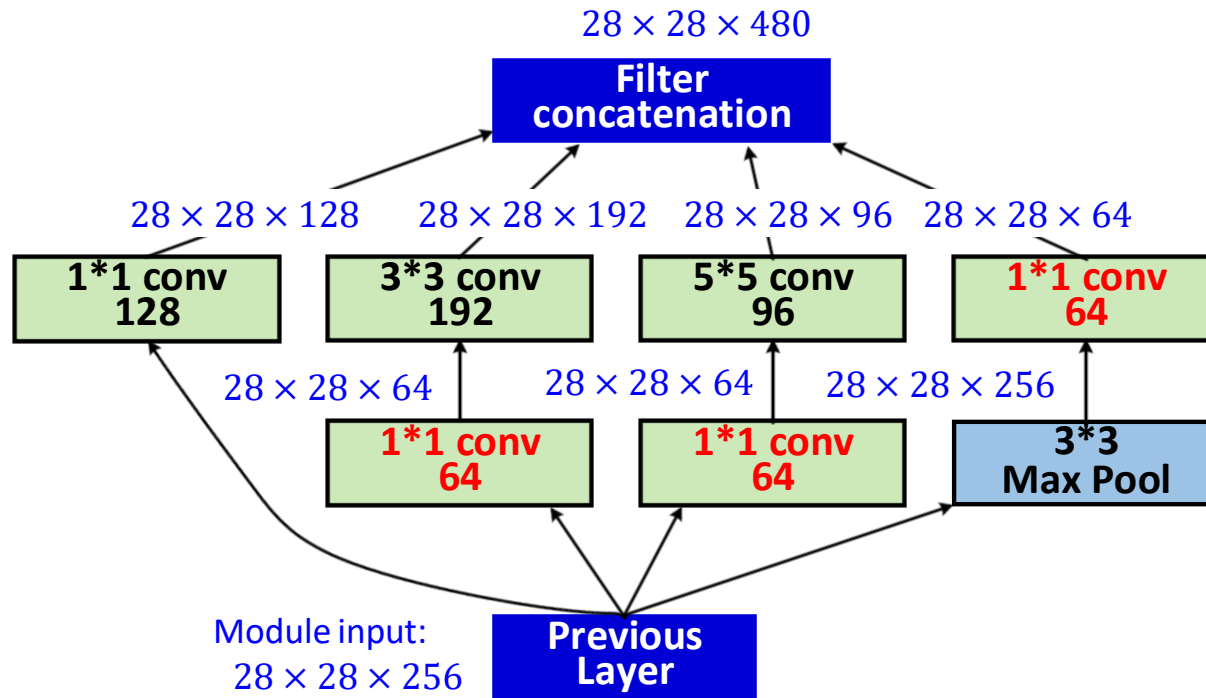


Preserves spatial dimensions, reduces depth!

Projects depth to lower dimension (combination of feature maps)

# Case Study: GoogLeNet

[Szegedy et al., 2014]



Inception module with dimension reduction

Using same parallel layers as naive example, and adding “ $1 \times 1$  conv, 64 filter” bottlenecks:


## Conv Ops:

$[1 \times 1 \text{ conv}, 64] 28 \times 28 \times 64 \times 1 \times 1 \times 256$   
 $[1 \times 1 \text{ conv}, 64] 28 \times 28 \times 64 \times 1 \times 1 \times 256$   
 $[1 \times 1 \text{ conv}, 128] 28 \times 28 \times 128 \times 1 \times 1 \times 256$   
 $[3 \times 3 \text{ conv}, 192] 28 \times 28 \times 192 \times 3 \times 3 \times 64$   
 $[5 \times 5 \text{ conv}, 96] 28 \times 28 \times 96 \times 5 \times 5 \times 64$   
 $[1 \times 1 \text{ conv}, 64] 28 \times 28 \times 64 \times 1 \times 1 \times 256$

**Total: 358M ops**

Compared to 854M ops for naive version Bottleneck can also reduce depth after pooling layer



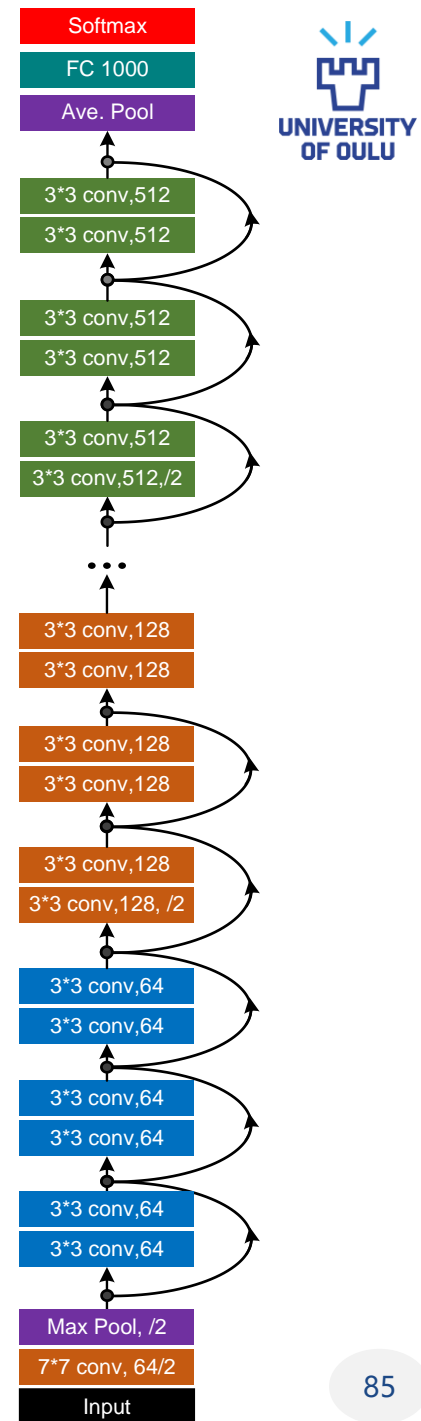



# Case Study: ResNet

[He et al., 2015]

Very deep networks using residual connections


- 152-layer model for ImageNet
- ILSVRC'15 classification winner (3.57% top 5 error)
- Swept all classification and detection competitions in ILSVRC'15 and COCO'15!



# Case Study: ResNet

[He et al., 2015]

What happens when we continue stacking deeper layers on a “plain” convolutional neural network?



Q: What’s strange about these training and test curves?

[Hint: look at the order of the curves]

56-layer model performs worse on both training and test error.

→The deeper model performs worse, but it’s not caused by overfitting!

# Case Study: ResNet

[He et al., 2015]


Hypothesis: the problem is an *optimization* problem, deeper models are harder to optimize

- The deeper model should be able to perform at least as well as the shallower model.
- A solution by construction is copying the learned layers from the shallower model and setting additional layers to identity mapping.


# Case Study: ResNet

[He et al., 2015]

Solution: Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



$$H(x) = F(x) + x$$



Use layers to fit residual  
 $F(x) = H(x) - x$   
 instead of  $H(x)$  directly




# Case Study: ResNet

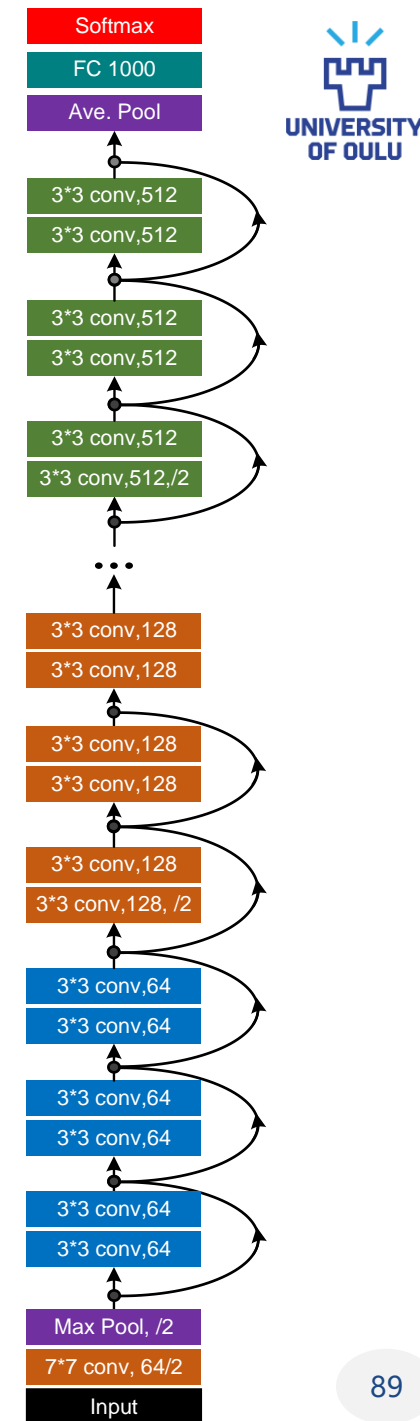
[He et al., 2015]

Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double the number of filters and subsample spatially using stride 2
- Additional conv layer at the beginning
- No FC layers at the end (only FC 1000 to output classes)



Total depths of 34, 50, 101, or 152 layers for ImageNet



# Case Study: ResNet

[He et al., 2015]

For deeper networks (ResNet50+), use “bottleneck” layer to improve efficiency (similar to GoogLeNet)

ILSVRC 2015 classification winner (3.6% top 5 error)→better than “human performance”!


## Experimental Results

- Able to train very deep networks without degrading (152 layers on ImageNet, 1202 on Cifar)
- Deeper networks now achieve lowering training error as expected
- Swept 1st place in all ILSVRC and COCO 2015 competitions

1x1 conv, 256 filters projects back to 256 feature maps (28x28x256)

3x3 conv operates over only 64 feature maps

1x1 conv, 64 filters to project to 28x28x64



# Typical CNN Architectures

## Case Studies

- AlexNet
- VGG
- GoogLeNet
- ResNet

about other architecture...

- NIN (Network in Network)
- Wide ResNet
- ResNeXT
- Stochastic Depth
- DenseNet
- SENet
- FractalNet
- SqueezeNet

No time to talk about, leave for yourself

# Summary: CNN Architectures

- VGG, GoogLeNet, ResNet all in wide use, available in model zoos
- ResNet current best default
- Trend towards extremely deep networks
- Significant research centers around design of layer / skip connections and improving gradient flow
- Even more recent trend towards examining necessity of depth vs. width and residual connections
- Lots of research in network compression and acceleration
- Lots of research in network architecture search

Next time: Training Tips for DNNs